

Gennoor Tech

TRAIN. INNOVATE. BUILD.

GENNOOR ACADEMY · SELF-PACED ·
40 MIN

Building Trusted Support Bots

A ~40-minute course for CS leaders, Support Ops, and CX Ops — containment math, knowledge grounding, escalation triggers, conversational design, voice, quality monitoring, 2-quarter rollout. Grounded in 2024–26 cases.

8

CHAPTERS

40 min

DURATION

Intermediate

LEVEL

FREE

CERTIFICATE

Course Overview

For CS leaders and CX Ops past basic AI literacy. Eight chapters across the support-bot surface: the CS AI landscape (Zendesk CX Trends 2026 84% strategic priority, Klarna up-then-down, Salesforce Agentforce 36% F100); containment vs deflection vs resolution math (McKinsey cost stack \$9–14 human vs \$1.50–3 autonomous); knowledge grounding (Air Canada Moffatt precedent, NYC MyCity, McKinsey 1-in-12 vs 1-in-200+ with retrieval); the 5 escalation triggers (Gallup CSAT-decline #1 reason, Chevy Tahoe \$1, warm-handoff design); conversational design + disclosure (EU AI Act Art. 50 Aug 2 2026, Pew 70% prefer disclosure, DPD haiku, 6 design rules + 3-sentence pattern); voice support AI (Talkdesk 25% inbound, Cognigy 250+ F2000, assist vs autonomous architectures, <900ms latency); quality monitoring (Forrester #3 failure, Salesforce 27-point gap, 5 metrics + shadow QA); and the close — 2-quarter rollout playbook with interactive builder, 4 trust trip-wires, and the through-line: AI for volume, humans for value.

FORMAT: Self-paced · 8 chapters · audio-narrated slides · interactive end-of-course builder · free certificate on completion

Who Is This For?

1

CX leaders

2

Customer service VPs

3

Support Operations
heads

4

Contact centre
leadership

Learning Outcomes

- 1 The honest CS-AI landscape — Zendesk CX Trends 84%, Klarna up-then-down lesson, Salesforce Agentforce reality
- 2 Containment vs deflection vs resolution math — the cost stack from McKinsey (\$9-14 human vs \$1.50-3 autonomous)
- 3 Knowledge grounding — Air Canada Moffatt precedent, NYC MyCity lesson, the retrieval-or-don't-build rule
- 4 The 5 escalation triggers — Gallup CSAT-decline #1, warm-handoff design, the Chevy Tahoe \$1 lesson
- 5 Conversational design + disclosure — EU AI Act Article 50 (Aug 2 2026), DPD haiku case, 6 design rules
- 6 Voice support AI — Talkdesk 25% inbound, sub-900ms latency, assist vs autonomous architectures
- 7 Quality monitoring — Forrester #3 failure mode, the 27-point gap, shadow QA + 5 metrics
- 8 A 2-quarter rollout with 4 trust trip-wires and the through-line: AI for volume, humans for value

8-Chapter Curriculum

Self-paced, audio-narrated slide format. Each chapter ends with a "Monday move" the learner can apply that week.

01

The CS AI landscape

CHAPTER 1 · 6 MIN · NARRATED + INTERACTIVE

- Zendesk CX Trends 2026 (84% strategic priority, 71% deployed).
- Avg containment 38–52% vs top-decile 70%+.
- Klarna up-then-down (2024 !' 2025).
- Salesforce, Intercom, Zendesk, Cognigy, Talkdesk landscape. 4 failure cases: Air Canada, DPD, NYC MyCity, Chevy Tahoe \$1.

02

Containment vs deflection · the math

CHAPTER 2 · 5 MIN · NARRATED + INTERACTIVE

- Containment "`` deflection "`` resolution — only resolution matters to CSAT.
- McKinsey cost stack: \$9–14 human vs \$1.50–3 autonomous.
- The whiteboard formula.
- Klarna reframed in money: appropriate containment, not maximum containment.

03

Knowledge grounding

CHAPTER 3 · 5 MIN · NARRATED + INTERACTIVE

- Moffatt v.
- Air Canada — "Air Canada is responsible for all information on its website, including from a chatbot." NYC MyCity harm.
- McKinsey 1-in-12 vs 1-in-200+.
- The grounding stack: retrieval, citations, confidence thresholds.
- The stale-KB trap.

04

Escalation triggers

CHAPTER 4 · 5 MIN · NARRATED + INTERACTIVE

- Gallup #1 CSAT-decline reason: "the bot kept me in a loop." Chevy Tahoe \$1 case.
- The 5 triggers: low confidence, sensitive intent, repeated misunderstanding, regulated topic, user request.
- The 4-rule warm handoff.
- Forrester #1 failure reason.

05

Conversational design + disclosure

CHAPTER 5 · 5 MIN · NARRATED + INTERACTIVE

- EU AI Act Article 50 from Aug 2 2026 — transparency obligation.
- Pew Research: 70% of consumers prefer disclosure.
- The three-sentence disclosure pattern.
- DPD haiku case.
- The 6 design rules: short, no apologies-as-deflection, named, single-purpose, on-rails, no fake humanness.

06

Voice support AI

CHAPTER 6 · 5 MIN · NARRATED + INTERACTIVE

- Talkdesk 25% inbound + 18% AHT down.
- Cognigy 250+ F2000. \$4/call autonomous vs \$6.50 human.
- Two architectures: agent-assist vs autonomous voice agent.
- Latency <900ms + barge-in.
- Build the architecture-per-intent matrix.

07

Quality monitoring

CHAPTER 7 · 5 MIN · NARRATED + INTERACTIVE

- Forrester #3 failure reason: no QA.
- Salesforce State of Service: 27-point CX gap between teams with active monitoring vs without.
- The 5 metrics: misrouted escalation rate, CSAT on AI-touched, repeat contact, transitions/session, abandonment.
- Shadow QA at 2% sample.

08

Making it stick: your CS AI roadmap

CHAPTER 8 · 4 MIN · NARRATED + INTERACTIVE

- 2 use cases, 2 quarters, 1 quality bar.
- Q1 = FAQ deflection grounded + all 5 triggers + shadow QA.
- Q2 = one transactional intent. 4 trust trip-wires not to cross.
- Interactive 2-quarter roadmap builder included.
- The through-line: AI for volume, humans for value.