

Gennoor Tech

TRAIN. INNOVATE. BUILD.

GENNOOR ACADEMY · SELF-PACED ·
45 MIN

Evaluating AI Output

A 45-minute discipline for professionals reviewing AI-generated work — spotting hallucinations, checking sources, building a verification habit.

7

CHAPTERS

45 min

DURATION

Intermediate

LEVEL

FREE

CERTIFICATE

Course Overview

For ICs and managers reviewing AI-generated work — their own, their team's, or a vendor's. Seven chapters: the fluency illusion (Microsoft FAccT 2024, Anthropic sycophancy, Stanford HAI 2026); accuracy vs. usefulness as separate tests; three hallucination patterns (confident fabrication, plausible detail, stale fact) anchored in NIST AI 600-1 and the Vectara HHEM leaderboard; citation evaluation with the Mata v. Avianca anchor, 1,353+ court cases, Sixth Circuit \$30K sanctions, and Deloitte Australia's AU\$290K refund; demographic and regional bias with named cases (Bloomberg study, EEOC iTutorGroup, Workday Mobley class action); the verification habit grounded in Lally's 66-day study and BJ Fogg's B=MAP formula; and the close — your one-page verification playbook with three never-skip checks, one escalation rule, and the Friday review.

FORMAT: Self-paced · 8 chapters · audio-narrated slides · interactive end-of-course builder · free certificate on completion

Who Is This For?

1

Professionals reviewing AI work

2

Editors + reviewers

3

Quality + compliance teams

4

AI-adopting knowledge workers

Learning Outcomes

- 1 How to spot hallucinations fast — the textual signals AI invents vs verifies
- 2 How to check sources in 30 seconds — the verification routine that scales
- 3 The 3-pass review pattern — facts, voice, audience
- 4 How to build a quality bar that your team can apply uniformly — not just the senior reviewer
- 5 When to ask the AI to critique its own output — and when that's circular
- 6 The verification habit pattern — paired with one tool, daily, until it's muscle memory
- 7 A 2-week practice plan that turns "AI output" into "trusted draft"

7-Chapter Curriculum

Self-paced, audio-narrated slide format. Each chapter ends with a "Monday move" the learner can apply that week.

01

Why AI evaluation is harder than it looks

CHAPTER 1 · 6 MIN · NARRATED + INTERACTIVE

- The fluency illusion is real.
- Microsoft FAccT 2024 study (404 participants), Anthropic sycophancy paper, Stanford HAI 2026 — confident writing convinces us regardless of accuracy.

02

Accuracy vs. usefulness

CHAPTER 2 · 6 MIN · NARRATED + INTERACTIVE

- Two different tests.
- Accurate-and-useless.
- Useful-and-wrong.
- The rule: whichever failure is hardest to recover from goes first.
- The Deloitte Australia anchor.

03

Spotting hallucinations in 3 patterns

CHAPTER 3 · 7 MIN · NARRATED + INTERACTIVE

- Confident fabrication.
- Plausible detail.
- Stale fact.
- NIST taxonomy, OpenAI "why models hallucinate", Vectara HHEM — reasoning models actually hallucinate MORE on long-form text.

04

Evaluating sources and citations

CHAPTER 4 · 7 MIN · NARRATED + INTERACTIVE

- 1,353+ court cases through 2026.
- Mata v.
- Avianca to the Sixth Circuit \$30K sanctions.
- Deloitte AU\$290K refund.
- The Nature 72% fake-citation finding.
- The 3-step citation check.

05

Spotting bias in outputs

CHAPTER 5 · 7 MIN · NARRATED + INTERACTIVE

- Bloomberg resume study (11% top-rank for Black women), EEOC iTutorGroup \$365K settlement, Workday Mobley class action, MMLU-ProX 30-point Swahili gap.
- Demographic + regional patterns.

06

Building your verification habit

CHAPTER 6 · 6 MIN · NARRATED + INTERACTIVE

- Lally's 66-day median to automaticity.
- Moore's overprecision research.
- BJ Fogg's B=MAP formula.
- The 5-minute three-step routine that survives week three.

07

Making it stick: your verification playbook

CHAPTER 7 · 5 MIN · NARRATED + INTERACTIVE

- Three never-skip checks.
- One escalation rule.
- The Friday review.
- Fill in the playbook builder, download the markdown, print it, pin it where you can see it.