

Gennoor Tech

TRAIN. INNOVATE. BUILD.

GENNOOR ACADEMY · SELF-PACED ·
36 MIN

MLOps for LLMs

A ~36-minute advanced course for ML and platform engineers. 5 LLM-vs-classical-ML differences, 3 drift types, 4 cost patterns, 60-sec rollback discipline, 4 incident types, the eval-case-it post-mortem. Voice: Andrew.

8

MODULES

36 min

DURATION

Advanced

LEVEL

FREE

CERTIFICATE

Course Brief

DURATION	36 min · 8 modules · self-paced
FORMAT	Audio-narrated slide chapters with synchronized progress tracking. Interactive end-of-course Markdown builder you take to the next meeting. Free verifiable certificate on completion.
AUDIENCE	Business + technology leaders, function heads + transformation owners, L&D + internal trainers, AI-curious managers.
PREREQUISITES	No technical prerequisites. The material is written for business decision-makers and assumes only general business literacy.

Course Outcomes

- 1 Five differences that make LLMOps its own discipline — non-determinism, generative outputs, third-party black boxes, per-request cost variability, adversarial users
- 2 3-tier production evaluation (light + targeted + gold sampling), LLM-as-judge done right (different family, structured rubric, human on disagreement), PII discipline
- 3 Three distinct drift types — input, output, behavioral — with regression-set replay to catch silent vendor model updates
- 4 Four cost optimization patterns (compress + cache + route + smaller-plus-retrieval) and three anti-patterns that break quality or correctness
- 5 One-artifact versioning (prompt + model + eval) with shadow/canary/feature-flag rollouts and 60-second rollback as a tested capability
- 6 OpenTelemetry GenAI semantic conventions + MLflow vs LangSmith for trace search, replay, and eval linking
- 7 Four incident types (quality, drift, abuse, vendor outage) with 5-question triage and the eval-case-it post-mortem that produces durable regression tests
- 8 5-section LLMOps runbook + 4 trust trip-wires + interactive runbook builder for the platform team

Table of Content

8 modules · 36 min of audio-narrated learning · each module ends with a Monday-move you can apply that week.

MODULE 01

LLMOps vs MLOps

~5 MIN · NARRATED + INTERACTIVE

- 5 differences (non-determinism
 - generative
 - third-party
 - cost variability
 - adversarial)
 - what carries
-

MODULE 02

Model evaluation in production

~5 MIN · NARRATED + INTERACTIVE

- 3-tier sampling (light
 - targeted
 - gold)
 - LLM-as-judge done right
 - PII discipline for regulated contexts.
-

MODULE 03

Drift detection

~5 MIN · NARRATED + INTERACTIVE

- 3 drift types (input
 - output
 - behavioral)
 - regression-set replay
 - actionable vs cosmetic alerts.
-

MODULE 04**Cost optimization patterns**

~5 MIN · NARRATED + INTERACTIVE

- The cost surprise
- 4 patterns that work (compress
- cache
- route
- smaller+retrieval)
- 3 anti-patterns

MODULE 05**Versioning prompts and models together**

~5 MIN · NARRATED + INTERACTIVE

- One-artifact discipline (prompt + model + eval triple)
- 3 rollout patterns (shadow
- canary
- feature flag)
- 60-sec rollback.

MODULE 06**Observability with traces**

~5 MIN · NARRATED + INTERACTIVE

- Why traces matter more
- OpenTelemetry GenAI
- Tooling stack
- A/B testing on traces

MODULE 07**Incident response for LLM failures**

~5 MIN · NARRATED + INTERACTIVE

- 4 incident types (quality
- drift
- abuse
- vendor outage)
- 5-question triage
- the eval-case-it post-mortem.

MODULE 08**Capstone — your LLMOps runbook**

~6 MIN · NARRATED + INTERACTIVE · CAPSTONE BUILDER

- The principle
- 5-section runbook
- 4 trust trip-wires
- Interactive · runbook builder

Taking This Further

The course is the starting point. The same content powers three concrete next steps you can engage Gennoor Tech to deliver inside your organisation:

4-week Pilot

- Fixed scope · fixed price · production artifact
- A working AI capability tied to the content of this course — set up on your data and tooling
- Suitable for: a single function or business unit ready to prove value

Org-wide Rollout

- 90-day adoption programme · executive + frontline tracks
- Champions network, prompt libraries, governance pack, weekly office hours
- Suitable for: 200-5,000-seat rollouts of Copilot, ChatGPT Enterprise, or similar

Continuous Build

- Quarterly retainer · embedded engineering + delivery + governance
- Multiple parallel AI initiatives sustained against an evolving roadmap
- Suitable for: large enterprises with portfolio-scale AI ambitions

Get in touch

Tell us about your AI ambitions, your team, and the constraints you are working under. We respond within one business day across India, GCC, SEA, US, and EU time zones.

enquiry@gennoor.com · www.gennoor.com · +91 · +971 · +1