

Gennoor Tech

TRAIN. INNOVATE. BUILD.

GENNOOR ACADEMY · SELF-PACED ·
36 MIN

Open-Source LLMs for Enterprise

A ~36-minute decision guide for senior engineers and tech directors evaluating or operating open-source LLMs at enterprise scale. Voice: Andrew.

8

MODULES

36 min

DURATION

Intermediate

LEVEL

FREE

CERTIFICATE

Course Brief

DURATION	36 min · 8 modules · self-paced
FORMAT	Audio-narrated slide chapters with synchronized progress tracking. Interactive end-of-course Markdown builder you take to the next meeting. Free verifiable certificate on completion.
AUDIENCE	Business + technology leaders, function heads + transformation owners, L&D + internal trainers, AI-curious managers.
PREREQUISITES	No technical prerequisites. The material is written for business decision-makers and assumes only general business literacy.

Course Outcomes

- 1 Open-source isn't free — it's a different cost curve. Three real reasons (sovereignty + residency, cost at scale, roadmap control) and 3 cases hosted API still wins
- 2 Four model families compared (Llama, Mistral, Phi, Qwen) with the license traps that catch teams and the size discipline that resists "bigger is better"
- 3 Self-hosting with Ollama vs vLLM vs TGI vs Azure ML — the gap between demo-grade benchmarks and production-grade concurrency
- 4 Fine-tuning approaches (LoRA · QLoRA · full) and the "try retrieval first" discipline — when fine-tuning is the wrong shortcut
- 5 5-component TCO with honest breakeven math — real breakeven is 20-50M tokens/day, not the headline GPU price
- 6 Three sovereign deployment tiers and what regulators + CISOs actually accept across India, GCC, US, EU
- 7 Operational discipline across model upgrades, security patches, and 3-layer observability (GPU · model · request)
- 8 A 4-section decision pack + 4 trust trip-wires + interactive decision-pack builder for the steering committee

Table of Content

8 modules · 36 min of audio-narrated learning · each module ends with a Monday-move you can apply that week.

MODULE 01

Why open-source and when not

~5 MIN · NARRATED + INTERACTIVE

- The principle
 - 3 reasons for open-source
 - 3 cases hosted API wins
 - The honest framing
-

MODULE 02

Model selection

~5 MIN · NARRATED + INTERACTIVE

- 4 families (Llama
 - Mistral
 - Phi
 - Qwen)
 - 3 license traps
 - size discipline
-

MODULE 03

Self-hosting

~5 MIN · NARRATED + INTERACTIVE

- 4 serving stacks
 - production-grade vs demo-grade
 - vLLM sizing (memory
 - quant
 - concurrency
 - P99).
-

MODULE 04**Fine-tuning approaches**

~5 MIN · NARRATED + INTERACTIVE

- Don't-fine-tune trap (retrieval first)
- 3 approaches (LoRA
- QLoRA
- full)
- when each fits
- when not at all.

MODULE 05**Cost and performance tradeoffs**

~5 MIN · NARRATED + INTERACTIVE

- 5-component TCO
- Breakeven math
- Concurrency reality
- Honest scoring · 2 questions

MODULE 06**Air-gapped and sovereign deployment**

~5 MIN · NARRATED + INTERACTIVE

- 3 deployment tiers
- 4 architectural requirements
- regional regulator acceptance (India
- GCC
- US
- EU).

MODULE 07**Operational considerations**

~5 MIN · NARRATED + INTERACTIVE

- 3 upgrade types (model
- serving stack
- quantization)
- 3 observability layers (GPU
- model
- request)

MODULE 08**Capstone — OSS LLM decision pack**

~6 MIN · NARRATED + INTERACTIVE · CAPSTONE BUILDER

- The principle
- 4-section decision pack
- 4 trust trip-wires
- Interactive · decision pack builder

Taking This Further

The course is the starting point. The same content powers three concrete next steps you can engage Gennoor Tech to deliver inside your organisation:

4-week Pilot

- Fixed scope · fixed price · production artifact
- A working AI capability tied to the content of this course — set up on your data and tooling
- Suitable for: a single function or business unit ready to prove value

Org-wide Rollout

- 90-day adoption programme · executive + frontline tracks
- Champions network, prompt libraries, governance pack, weekly office hours
- Suitable for: 200-5,000-seat rollouts of Copilot, ChatGPT Enterprise, or similar

Continuous Build

- Quarterly retainer · embedded engineering + delivery + governance
- Multiple parallel AI initiatives sustained against an evolving roadmap
- Suitable for: large enterprises with portfolio-scale AI ambitions

Get in touch

Tell us about your AI ambitions, your team, and the constraints you are working under. We respond within one business day across India, GCC, SEA, US, and EU time zones.

enquiry@gennoor.com · www.gennoor.com · +91 · +971 · +1