

**Gennoor Tech**

TRAIN. INNOVATE. BUILD.

GENNOOR ACADEMY · SELF-PACED ·  
55 MIN

# Prompting Mastery

A 55-minute deep dive for power users and developers — structured outputs, chain-of-thought, few-shot, evaluation, and the prompt-library habits that compound.

**8**

CHAPTERS

**55 min**

DURATION

**Intermediate**

LEVEL

**FREE**

CERTIFICATE

## Course Overview

For analysts, developers, and senior individual contributors who already use ChatGPT, Claude, or Copilot every day and want to move from ad-hoc prompts to repeatable, evaluable patterns. Eight chapters: principles that separate amateur from professional prompts; reliable JSON via structured outputs; chain-of-thought where it earns its keep; few-shot patterns that generalise; eval-driven iteration instead of vibe-checking; library structure, ownership, and versioning; the adversarial / prompt-injection chapter with real 2025 CVEs (EchoLeak, CamoLeak, Slack AI); and the close — the four things that decide whether the library lives past month three.

**FORMAT:** Self-paced · 8 chapters · audio-narrated slides · interactive end-of-course builder · free certificate on completion

## Who Is This For?

**1**

Power users +  
developers

**2**

AI engineering teams

**3**

Knowledge workers  
going deeper

**4**

Internal AI champions

## Learning Outcomes

- 1 Structured outputs that survive downstream parsing — JSON, XML, tables, citations
- 2 Chain-of-thought prompting — when it helps, when it harms, how to constrain it
- 3 Few-shot patterns that lift accuracy without bloating tokens
- 4 How to build and curate a prompt library — the habits that compound over months
- 5 Evaluation framework — how to know your prompt is good without running 1000 trials manually
- 6 Cost + latency awareness — when to use which model, the 3-axis tradeoff (quality, cost, speed)
- 7 Power-user tactics — temperature, system prompts, response shaping, error recovery

## 8-Chapter Curriculum

Self-paced, audio-narrated slide format. Each chapter ends with a "Monday move" the learner can apply that week.

01

### Prompting principles

CHAPTER 1 · 8 MIN · NARRATED + INTERACTIVE

- The five things that separate amateur prompts from professional ones — and the chat-vs-reasoning model toggle that changes which principles dominate.

02

### Structured outputs (JSON)

CHAPTER 2 · 7 MIN · NARRATED + INTERACTIVE

- Stop parsing free text.
- Build a schema, force JSON Schema strict mode, and watch the model return parseable JSON on the first try every time.

03

### Chain-of-thought reasoning

CHAPTER 3 · 7 MIN · NARRATED + INTERACTIVE

- A 2x2 of when chain-of-thought earns its latency cost — and when adding "think step by step" makes outputs worse, not better.

04

### Few-shot patterns

CHAPTER 4 · 7 MIN · NARRATED + INTERACTIVE

- Examples that generalise vs. examples that memorise — a similarity scorer that shows the precise distance between your few-shots and your actual input distribution.

05

### Evaluating outputs

CHAPTER 5 · 6 MIN · NARRATED + INTERACTIVE

- A mini eval runner you can drop into CI tomorrow — happy path, edge cases, adversarial inputs — that turns "looks good to me" into a measured pass rate.

06

### Prompt libraries and reuse

CHAPTER 6 · 5 MIN · NARRATED + INTERACTIVE

- Four files per prompt — prompt, schema, evals, changelog — plus owner, versioning, and CODEOWNERS.
- The structure that survives turnover.

07

### Edge cases and adversarial prompts

CHAPTER 7 · 8 MIN · NARRATED + INTERACTIVE

- EchoLeak (CVE-2025-32711), CamoLeak, and Slack AI — the three 2025 incidents that prove prompt injection is now a production problem.
- The four defences that actually work.

08

## Making your library stick

CHAPTER 8 · 7 MIN · NARRATED + INTERACTIVE

- Owner, weekly review, evals in CI, frictionless contribution — the four things that decide whether the library lives past month three.
- The 4-week rollout plan, downloadable per team size.